Markwise

Pre-publication brand governance that explains, cites and suggests, and leaves the decision to people.

Round 1 Submission: The Plan
Problem statement: Inkloom AI, Intelligent Brand-Governance Platform

Team: arvindshekhawat845 | Members: Arvind Shekhawat, Alok Kumar Singh, Niket Malik, Kapil Chaudhary

In one paragraph

Brand mistakes rarely come from bad intent. They come from many people, agencies and AI tools publishing quickly against guidelines that live in a PDF nobody opens. Markwise sits between 'draft' and 'publish'. It checks content against the organisation's own approved guidelines, fact sheet, glossary and asset library, and every finding names the exact source rule it came from. It proposes a fix that keeps what the author meant. A human always makes the final call, and the system learns where the guidelines themselves are unclear.

Contents

1 Problem 2 Users and setting 3 Evidence 4 Solution and workflow 5 Features and the three improvements 6 Technology 7 Responsible design 8 Plan, roles, impact, risks

1. The problem

The more people and tools an organisation publishes through, the less any one person can see everything that goes out. Guidelines are static documents, facts change, and generative AI now lets anyone produce polished copy in seconds without knowing the brand rules. The damage is often noticed by customers or competitors before the brand team.

Four failure patterns we are designing against

Pattern	Example	Why current tools miss it
Unsupported claim	'Cuts onboarding time by 60%' when the approved figure is 40% from one study	Spell and grammar checkers do not know the company's facts
Outdated information	Old pricing, discontinued product names, previous legal line	No link between content and a versioned source of truth
Voice and terminology drift	'Customers' in one team, 'clients' in another; a casual tone in a regulated product	Guideline PDFs are advice, not enforced checks
Visual and accessibility errors	Off-palette colour, distorted logo, low contrast, missing alt text	Reviewed manually, late, and inconsistently

Why existing options fall short

- **Style checkers** (grammar tools with style guides) check language, but not facts, claims or visuals, and rarely cite why.
- **Digital asset managers** store approved assets but do not review a finished post, document or slide.
- **Manual brand review** does not scale, becomes a bottleneck, and gives feedback like 'this feels off-brand' without a rule behind it.
- **General chatbots** can critique copy but have never seen the company's guidelines, hallucinate rules, and offer no audit trail.

2. Target users and setting

User	Need	What Markwise gives them
Content creator (marketer, sales, agency writer)	Publish fast without fear of getting it wrong; know what to change and why	Instant pre-check with plain-language reasons and one-click suggested edits
Brand / content manager	Stop repeat errors; stop being the bottleneck	Only escalated items reach them; dashboard of recurring issues
Authorised reviewer (legal, compliance, product)	Verify claims against approved facts	Claim-by-claim view with the linked source and an approve or reject action
Brand admin	Keep guidelines current and enforceable	Rule library with versions, and a list of rules that cause confusion

Setting. Mid-sized to large organisations (roughly 50 to 5,000 staff) that publish across web, social, email, decks and PDFs through in-house teams, agencies and AI writing tools. Regulated sectors (finance, health, education) have the highest stakes because claims carry legal weight.

3. Evidence that the problem is real

Our evidence has two parts: documented public cases that already exist, and a validation plan we will run with real users during the build window. We do not claim survey or interview results we have not yet collected.

Documented public cases

Case	What went wrong	Relevance
Air Canada chatbot (Moffatt v. Air Canada, 2024)	Customer-facing content gave wrong policy information; a tribunal held the company responsible	Unverified claims published without checking them against approved sources

Case	What went wrong	Relevance
CNET AI-written articles (2023)	AI-assisted finance articles were found to contain errors and were publicly corrected	AI-generated content published with no fact gate
Pepsi protest-themed ad (2017)	Ad was pulled within a day after backlash over tone and messaging	Tone and messaging problems found only after publication
Advertising regulators (ASCI, ASA, FTC)	Regular rulings against unsubstantiated or misleading promotional claims	Claim substantiation is a legal exposure, not just a style issue

Validation plan (to run in the build window)

Method	Target	Purpose
Interviews with marketers, content writers or agency staff	5 to 8	Learn how review happens today and where errors slip through
Short survey on brand errors caught after publishing	30+ responses	Measure how common post-publish corrections are
Test of a general AI chatbot on a flawed sample set we create	1 set	Show whether it misses claims or invents rules without our guidelines
User test of the Markwise prototype	3 to 5 users	Check findings are understood and suggestions keep the author's meaning

Illustrative example (our own test case). Approved fact: 'Reduces onboarding time by up to 40% (2025 customer study).' Draft post: 'We cut onboarding time by 60%, guaranteed!' Expected finding: *claim does not match approved fact FACT-014; 'guaranteed' is restricted by rule LEGAL-03*. Suggested edit: 'Customers reduced onboarding time by up to 40%.' This is the behaviour the prototype must deliver.

4. Proposed solution and workflow

Value proposition. Catch brand, fact, claim and accessibility problems *before* publication, with a reason and a source for every finding, in minutes instead of review-queue days, while keeping humans in charge of every final decision.

End-to-end workflow

#	Step	What happens
1	Set up the knowledge base	Admin imports guidelines (PDF/DOCX), approved facts, glossary and assets. The system structures them into versioned, testable rules and fact entries. The admin confirms them.
2	Submit content	Creator uploads text, an image, a PDF or a slide deck, and states the intent (channel, audience, goal).
3	Analyse	Deterministic checks run first (banned terms, palette, contrast, alt text, logo clear-space). Then an AI model checks tone, claims and facts using only retrieved approved sources.
4	Explain	Each finding shows: what, where in the content, why, the exact guideline or fact it cites, severity and a confidence level.
5	Suggest	A proposed rewrite keeps the creator's intent. It is shown as a tracked change that the creator can accept, edit or dismiss with a reason.
6	Review and approve	Clean content is auto-cleared to the creator's manager. High-risk items (claims, legal text) route to named reviewers. Comments, assignments, and versions are recorded.
7	Learn	Dismissals, overrides and repeat issues feed the insights dashboard and flag unclear guidelines.

5. Features and the three improvements we commit to

Core capabilities map directly to the brief: guideline and fact import, text/image/document review, tone, terminology, claim, visual and accessibility detection, explanations, suggestions, roles and approvals, insights, and confidentiality controls.

Three named improvements (we will report progress on these in Rounds 2 and 3)

Improvement	What it means	How we will measure it
1. Source-cited, abstaining findings	No finding without a linked rule or fact. If no approved source covers a claim, the system says 'cannot verify' instead of guessing.	Share of findings with valid citation (target 100%); hallucinated-rule rate on a test set
2. Intent-preserving suggestions	Rewrites are shown as a diff against the original, limited to the flagged span, with the option to explain or reject.	Accept rate; creators' rating of 'kept my meaning' in user tests
3. Guideline-gap insights	Detect rules that are often dismissed, contested or contradictory, and recurring error types per team or agency.	Number of ambiguous rules surfaced and fixed during testing

Scope discipline

Build window MVP	Stretch	Future
Text, image and PDF review; rules + AI checks; citations; roles; comments; versions; basic insights	DOCX and slides; tone scoring per channel; Slack/CMS plugin	Video/audio, multilingual brand rules, Figma plugin, SSO and on-prem model option

6. Technology, integrations and data

Layer	Choice	Reason
Frontend	React / Next.js, Tailwind	Fast to build; diff and comment views
Backend	Python (FastAPI)	Strong document, image and AI libraries

Layer	Choice	Reason
Database	PostgreSQL with pgvector; object storage for files	One store for rules, facts, versions, audit log and retrieval
AI	LLM via API for classification and rewriting, with retrieval over approved sources only; vision model for image checks	Grounded answers; model is swappable
Deterministic engine	Rules for banned terms, regex, colour palette distance, WCAG contrast ratio, alt text presence	Reliable, free, explainable, and not dependent on an AI model
Parsing	PDF and DOCX parsers, OCR for scanned or image text	Supports the multiple formats in the brief
Auth and roles	Role-based access (creator, manager, reviewer, admin)	Matches the approval workflow

Data sources. Only material the organisation uploads. For the demo we will use a fictional company's guidelines, fact sheet and assets that we author ourselves, plus public sample posts. No private data is required. Integrations beyond upload (Slack, CMS, Google Drive) are stretch goals through webhooks and connectors.

7. Responsible design

Concern	Our approach
Human control	AI only suggests. Nothing is auto-edited or auto-published. Every override is logged with a reason. High-risk categories always need a named human reviewer.
Explainability	Each finding: rule or fact ID, quoted source passage, severity, confidence. 'Cannot verify' is a valid outcome.
Confidentiality	Per-workspace data isolation, encryption in transit and at rest, role-based access, a setting to disable external AI calls, and a retention and deletion control.
Transparency	The UI discloses when AI or an external service analysed content, which provider, and what was sent. Content is not used to train models.
Fairness and quality	Tone checks judge against the organisation's own voice guide, not an assumed 'correct' style, and avoid penalising dialects or non-native phrasing without a rule.
Accessibility	Our own interface follows WCAG 2.1 AA: keyboard navigation, contrast, and labels.
Disclosure	All libraries, templates and AI-assisted code will be listed in the repository, as the rules require.

8. Development plan, roles, impact and risks

Phase	Focus	Output
Week 1	Knowledge base import, rule and fact schema, auth and roles	Admin can load guidelines and facts
Week 2	Text and image review engine, citations, explanations	End-to-end review of one asset
Week 3	Suggestions as diffs, approval workflow, comments, version history	Full creator to approver flow
Week 4	Insights dashboard, PDF support, user testing, hardening	Tested MVP and demo

Member	Role	Owns
Arvind Shekhawat	Product and AI lead	Prompting, retrieval, rule and fact schema, evaluation set
Alok Kumar Singh	Backend lead	API, database, rules engine, document parsers, approval workflow
Niket Malik	Frontend lead	Review screen, diff view, comments, dashboard, accessibility
Kapil Chaudhary	Research and QA	User interviews, survey, test cases, documentation, demo

Expected impact

- Fewer post-publication corrections and less review time per asset (we will measure against a manual baseline in user tests).
- Brand managers spend time on judgement calls instead of repeat fixes.
- Guideline owners get evidence about which rules to rewrite.

Key risks and mitigations

- **AI errors:** grounded retrieval, abstain option, deterministic checks first, and a labelled test set to track precision.
- **Scope creep:** text, image and PDF only for the MVP; others are stretch goals.
- **Guideline quality:** messy source documents are structured with admin confirmation instead of blind trust.